

Problem 4 (5 points)

Now you will try support vector classification on data with nonlinear decision boundaries. You will use the sklearn SVC tool on four datasets. Your job is to find an appropriate choice of kernel and regularization strength that does a qualitatively good job separating the data.

Run this cell first:

```
In [19]: import numpy as np
import matplotlib.pyplot as plt
from sklearn.svm import SVC
from matplotlib.colors import ListedColormap

# Plotting functions:
def plot_data(X,c,s=30):
    lims = [0,1]

    markers = [dict(marker="o", color="royalblue"), dict(marker="s", color="red"), dict(marker="x", color="green")]

    x,y = X[:,0], X[:,1]
    iter = 0
    for i in np.unique(c):
        marker = markers[iter]
        iter += 1
        plt.scatter(x[c==i], y[c==i], s=s, **(marker), edgecolor="black", linewidths=1)

def plot_SVs(svm, s=120):
    sv = svm.support_vectors_
    x, y = sv[:,0], sv[:,1]
    plt.scatter(x, y, s=s, edgecolor="black", facecolor="none", linewidths=1.5)

def plot_SV_decision_boundary(svm, margin=True, extend=True, shade_margins=False, shade_decision=True):
    ax = plt.gca()
    xlim = ax.get_xlim()
    ylim = ax.get_ylim()

    xrange = xlim[1] - xlim[0]
    yrange = ylim[1] - ylim[0]

    x = np.linspace(xlim[0] - extend*xrange, xlim[1] + extend*xrange, 200)
    y = np.linspace(ylim[0] - extend*yrange, ylim[1] + extend*yrange, 200)

    X,Y = np.meshgrid(x,y)
    xy = np.vstack([X.ravel(), Y.ravel()]).T
    P = svm.decision_function(xy)

    P = P.reshape(X.shape)
    ax.contour(X, Y, P, colors='k', levels=[0], linestyles=['-'])
    if margin:
        ax.contour(X, Y, P, colors='k', levels=[-1, 1], alpha=0.6, linestyles=['--'])

    if shade_margins:
        cmap = ListedColormap(["white", "lightgreen"])
        plt.pcolormesh(X,Y,np.abs(P)<1, shading="nearest", cmap=cmap, zorder=-999999)
```

```

if shade_decision:
    cmap = ListedColormap(["lightblue", "lightcoral"])
    pred = (svm.predict(xy).reshape(X.shape) == 1).astype(int)
    plt.pcolormesh(X, Y, pred, shading="nearest", cmap=cmap, zorder=-1000)

plt.xlim(xlim)
plt.ylim(ylim)

def plot(Xdata, ydata, svm_model=None, title=""):
    plt.figure(figsize=(5,5))
    plot_data(Xdata, ydata)
    if svm_model is not None:
        plot_SVs(svm_model)
        plot_SV_decision_boundary(svm_model, margin=True, shade_decision=True)
    plt.legend()
    plt.xlabel("$x_1$")
    plt.ylabel("$x_2$")
    plt.title(title)
    plt.show()

```


Loading the data



There are four datasets, all 2D and with X and y names as follows:


- X_a, y_a
- X_b, y_b
- X_c, y_c
- X_d, y_d

Run this cell to load and plot the data:

```
In [20]: xa0 = np.array([0.00806452, 0.0467742, -0.0145161, 0.0564516, 0.130645, 0.0887097, 0.2  
ya0 = np.array([0.418367, 0.540816, 0.565306, 0.704082, 0.720408, 0.822449, 0.883673,  
xa1 = np.array([0.324194, 0.304839, 0.372581, 0.369355, 0.453226, 0.546774, 0.724194,  
ya1 = np.array([0.540816, 0.320408, 0.25102, 0.185714, 0.120408, 0.0959184, 0.104082,  
Xa = np.concatenate([np.vstack([xa0,ya0]).T,np.vstack([xa1,ya1]).T],0)  
ya = np.array([-1,-1,-1,-1,-1,-1,-1,-1,-1,-1,-1,-1,-1,-1,-1,-1,-1,-1,-1,-1,-1,-1,-1])  
plot(Xa,ya,title="Dataset A")  
  
xb1 = np.array([0.43599,0.54966,0.42037,0.20465,0.29965,0.62113,0.13458,0.18444,0.8539  
xb2 = np.array([0.025926,0.43532,0.33033,0.61927,0.26683,0.52914,0.51358,0.78534,0.494  
yb = np.array([1,-1,1,1,1,-1,1,1,-1,-1,-1,1,1,1,1,-1,-1,-1,1,-1,-1,-1,1,1,1,1,1,1,1,1,1,1,1,  
Xb = np.vstack([xb1,xb2]).T  
plot(Xb,yb,title="Dataset B")  
  
xc1 = np.array([0.05,0.08871,0.18226,0.18871,0.27581,0.25323,0.10806,0.19516,0.25645,0.  
xc2 = np.array([0.58571,0.65918,0.63878,0.51633,0.51633,0.37755,0.3449,0.47551,0.69184  
Xc = np.vstack([xc1,xc2]).T  
yc = np.array([-1,-1,-1,-1,-1,-1,-1,-1,-1,-1,-1,1,1,1,1,1,1,1,1,1,1,1,1,1,1,1,1,1,1,1,1,1,1,  
plot(Xc,yc,title="Dataset C")  
  
xd1 = np.array([0.062903,0.08871,0.18548,0.33065,0.54032,0.68226,0.81129,0.91129,0.927  
xd2 = np.array([0.52041,0.67143,0.77347,0.88367,0.92857,0.83469,0.64286,0.54898,0.3898  
Xd = np.vstack([xd1,xd2]).T
```

[illegible]





Using the Kernel Trick


Now, train four SVC models, one for each dataset. Try out different combinations of 'kernel' and 'C', until you find a satisfactory classifier in each case.

Please generate a plot for each dataset showing the results of a trained support vector classifier, using the provided function:

```
plot(Xdata, ydata, svm_model, title)
```


```
In [21]: # YOUR CODE GOES HERE
# (Dataset A)

svm = SVC(kernel = "rbf", C = 10)
svm.fit(Xa, ya)
plot(Xa, ya, svm, title=f"C= 10 Kernel= rbf for Dataset A")
```



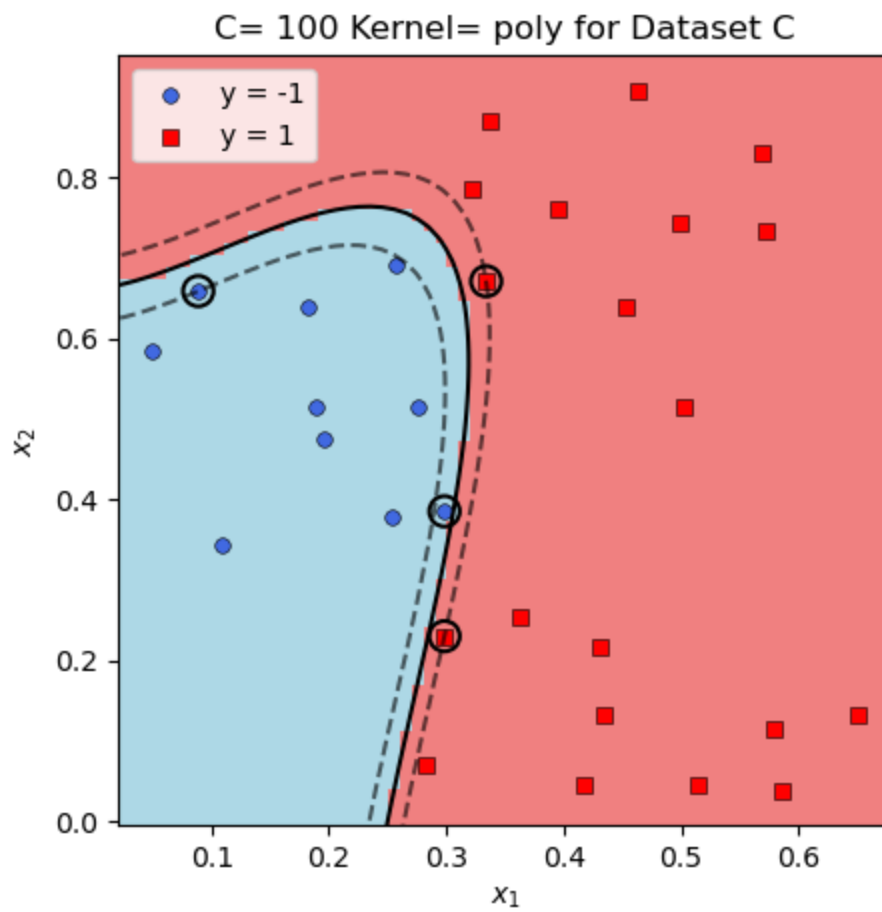
```
In [22]: # YOUR CODE GOES HERE
# (Dataset B)

svm = SVC(kernel = "rbf",C = 500)
svm.fit(Xb,yb)
plot(Xb,yb,svm,title=f"C= 500 Kernel= rbf for Dataset B")
```



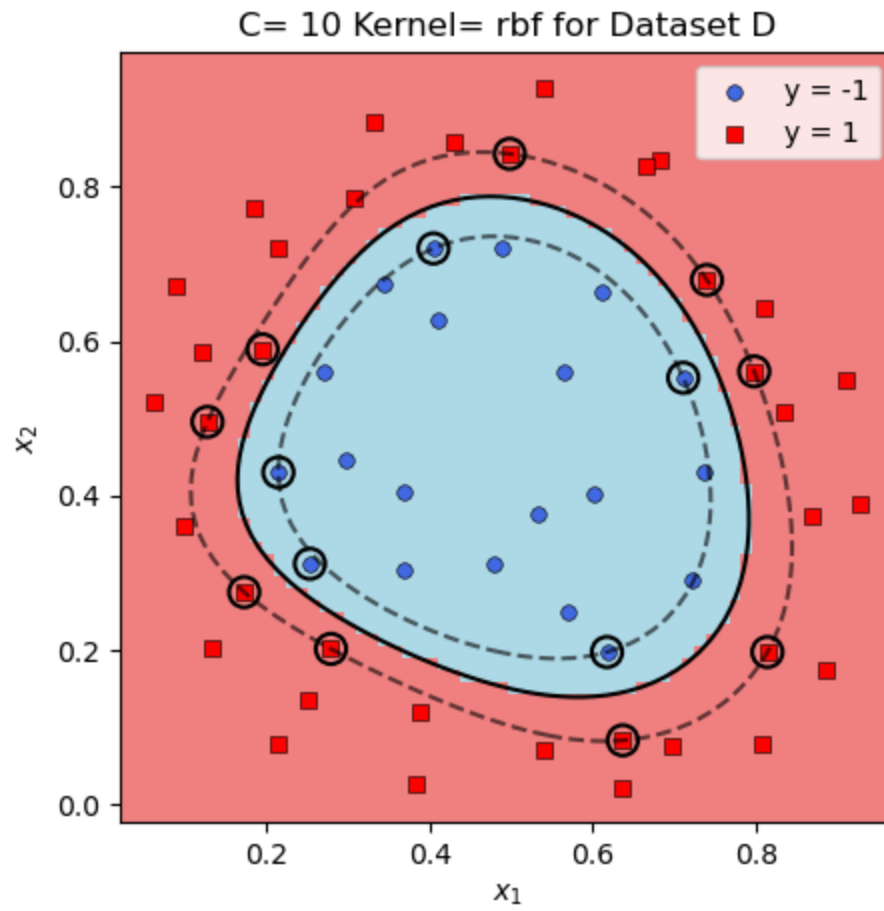
```
In [23]: # YOUR CODE GOES HERE
# (Dataset C)

svm = SVC(kernel = "poly",C = 100)
svm.fit(Xc,yc)
plot(Xc,yc,svm,title=f"C= 100 Kernel= poly for Dataset C")
```



```
In [24]: # YOUR CODE GOES HERE
# (Dataset D)

svm = SVC(kernel = "rbf",C = 10)
svm.fit(Xd,yd)
plot(Xd,yd,svm,title=f"C= 10 Kernel= rbf for Dataset D")
```

In []: