

M8-L2 Problem 2

Let's revisit the material phase prediction problem once again. You will use this problem to try multi-class classification in PyTorch. You will have to write code for a classification network and for training.

```
In [1]: import numpy as np
import matplotlib.pyplot as plt
from matplotlib.colors import ListedColormap
import torch
from torch import nn, optim
import torch.nn.functional as F

def plot_loss(train_loss, val_loss):
    plt.figure(figsize=(4,2),dpi=250)
    plt.plot(train_loss,label="Training")
    plt.plot(val_loss,label="Validation",linewidth=1)
    plt.legend()
    plt.xlabel("Epoch")
    plt.ylabel("Loss")
    plt.show()

def split_data(X, Y):
    np.random.seed(100)
    N = len(Y)
    train_mask = np.zeros(N, dtype=np.bool_)
    train_mask[np.random.permutation(N)[:int(N*0.8)]] = True
    train_x, val_x = torch.Tensor(X[train_mask,:]), torch.Tensor(X[np.logical_not(train_mask),:])
    train_y, val_y = torch.Tensor(Y[train_mask]), torch.Tensor(Y[np.logical_not(train_mask)])
    return train_x, val_x, train_y, val_y

x1 = np.array([7.4881350392732475,16.351893663724194,22.427633760716436,29.04883182996897,35.03654799338904,44.458941130
x2 = np.array([0.11120957227224215,0.1116933996874757,0.14437480785146242,0.11818202991034835,0.0859507900573786,0.09370
X = np.vstack([x1,x2]).T
y = np.array([0,2,2,2,2,2,0,2,2,2,2,2,0,0,2,0,1,2,0,0,1,1,1,2,0,1,0,1,1,1,0,0,1,1,1,1,])

X = torch.Tensor(X)
Y = torch.tensor(y,dtype=torch.long)

train_x, val_x, train_y, val_y = split_data(X,Y)
```

```

def plot_data(newfig=True):
    xlim = [0,52.5]
    ylim = [0,1.05]
    markers = [dict(marker="o", color="royalblue"), dict(marker="s", color="crimson"), dict(marker="D", color="limegreen")]
    labels = ["Solid", "Liquid", "Vapor"]

    if newfig:
        plt.figure(figsize=(6,4),dpi=250)

    x = X.detach().numpy()
    y = Y.detach().numpy().flatten()

    for i in range(1+max(y)):
        plt.scatter(x[y==i,0], x[y==i,1], s=40, **(markers[i]), edgecolor="black", linewidths=0.4,label=labels[i])


    plt.scatter(val_x[:,0], val_x[:,1],s=120,c="None",marker="o",edgecolors="black",label="Test point")

    plt.title("Phase of simulated material")
    plt.legend(loc="upper right")
    plt.xlim(xlim)
    plt.ylim(ylim)
    plt.xlabel("Temperature, K")
    plt.ylabel("Pressure, atm")
    plt.box(True)

def plot_model(model, res=200):
    xlim = [0,52.5]
    ylim = [0,1.05]
    xvals = np.linspace(*xlim,res)
    yvals = np.linspace(*ylim,res)
    x,y = np.meshgrid(xvals,yvals)
    XY = np.concatenate((x.reshape(-1,1),y.reshape(-1,1)),axis=1)
    XY = torch.Tensor(XY)
    color = model.predict(XY).reshape(res,res).detach().numpy()
    cmap = ListedColormap(["lightblue","lightcoral","palegreen"])
    plt.pcolor(x, y, color, shading="nearest", zorder=-1, cmap=cmap,vmin=0,vmax=2)
    return

plot_data()
plt.show()

```



Model definition

In the cell below, complete the definition for `PhaseNet`, a classification neural network.

- The network should take in 2 inputs and return 3 outputs.
- The network size and hidden layer activations are up to you.
- Make sure to use the proper activation function (for multi-class classification) at the final layer.
- The `predict()` method has been provided, to return the integer class value. You must finish `__init__()` and `forward()`.

```
In [2]: class PhaseNet(nn.Module):
    def __init__(self):
        super().__init__()
        # YOUR CODE GOES HERE
        N_in = 2
        N_hidden = 12
        N_out = 3
        self.lin1 = nn.Linear(N_in, N_hidden)
        self.lin2 = nn.Linear(N_hidden, N_hidden)
        self.lin3 = nn.Linear(N_hidden, N_hidden)
        self.lin4 = nn.Linear(N_hidden, N_hidden)
        #self.lin5 = nn.Linear(N_hidden, N_hidden)
        self.lin5 = nn.Linear(N_hidden, N_out)
        self.act = F.relu

    def predict(self, X):
        Y = self(X)
        return torch.argmax(Y, dim=1)

    def forward(self, X):
        # YOUR CODE GOES HERE
        X = self.lin1(X)
        X = self.act(X)
        X = self.lin2(X)
        X = self.act(X)
        X = self.lin3(X)
        X = self.act(X)
        X = self.lin4(X)
        X = self.act(X)
        X = self.lin5(X)
        #X = self.act(X)
        #X = self.lin6(X)
        X = F.softmax(X, dim = 1)

        return X
```

Training

Most of the training code has been provided below. Please add the following where indicated:

- Define a loss function (for multiclass classification)

- Define an optimizer and call it `opt`. You may choose which optimizer.

Make sure the training curves you get are reasonable.

```
In [3]: model = PhaseNet()

lr = 0.001
epochs = 1000

# Define loss function
# YOUR CODE GOES HERE
lossfun = nn.CrossEntropyLoss()

# Define an optimizer, `opt`
# YOUR CODE GOES HERE
opt = optim.Adam(params = model.parameters(), lr=lr)

train_hist = []
val_hist = []


for epoch in range(epochs+1):
    model.train()
    loss_train = lossfun(model(train_x), train_y)
    train_hist.append(loss_train.item())

    model.eval()
    loss_val = lossfun(model(val_x), val_y)
    val_hist.append(loss_val.item())

    opt.zero_grad()
    loss_train.backward()
    opt.step()
    if epoch % int(epochs / 25) == 0:
        print(f"Epoch {epoch:>4} of {epochs}:   Train Loss = {loss_train.item():.4f}   Validation Loss = {loss_val.item():.4f}")

plot_loss(train_hist, val_hist)
```


Epoch 0 of 1000:	Train Loss = 1.1018	Validation Loss = 1.0920
Epoch 40 of 1000:	Train Loss = 1.0839	Validation Loss = 1.0862
Epoch 80 of 1000:	Train Loss = 1.0454	Validation Loss = 1.0762
Epoch 120 of 1000:	Train Loss = 1.0017	Validation Loss = 1.0492
Epoch 160 of 1000:	Train Loss = 0.9548	Validation Loss = 1.0204
Epoch 200 of 1000:	Train Loss = 0.9200	Validation Loss = 0.9962
Epoch 240 of 1000:	Train Loss = 0.8903	Validation Loss = 0.9645
Epoch 280 of 1000:	Train Loss = 0.8470	Validation Loss = 0.9123
Epoch 320 of 1000:	Train Loss = 0.7806	Validation Loss = 0.8581
Epoch 360 of 1000:	Train Loss = 0.7065	Validation Loss = 0.8254
Epoch 400 of 1000:	Train Loss = 0.6772	Validation Loss = 0.7890
Epoch 440 of 1000:	Train Loss = 0.6631	Validation Loss = 0.7817
Epoch 480 of 1000:	Train Loss = 0.6544	Validation Loss = 0.7729
Epoch 520 of 1000:	Train Loss = 0.6482	Validation Loss = 0.7693
Epoch 560 of 1000:	Train Loss = 0.6433	Validation Loss = 0.7680
Epoch 600 of 1000:	Train Loss = 0.6390	Validation Loss = 0.7702
Epoch 640 of 1000:	Train Loss = 0.6357	Validation Loss = 0.7724
Epoch 680 of 1000:	Train Loss = 0.6330	Validation Loss = 0.7750
Epoch 720 of 1000:	Train Loss = 0.6310	Validation Loss = 0.7766
Epoch 760 of 1000:	Train Loss = 0.6294	Validation Loss = 0.7786
Epoch 800 of 1000:	Train Loss = 0.6281	Validation Loss = 0.7795
Epoch 840 of 1000:	Train Loss = 0.6272	Validation Loss = 0.7799
Epoch 880 of 1000:	Train Loss = 0.6265	Validation Loss = 0.7800
Epoch 920 of 1000:	Train Loss = 0.6259	Validation Loss = 0.7787
Epoch 960 of 1000:	Train Loss = 0.6255	Validation Loss = 0.7786
Epoch 1000 of 1000:	Train Loss = 0.6251	Validation Loss = 0.7786



Plot results

Plot your network predictions with the data by running the following cell. If your network has significant overfitting/underfitting, go back and retrain a new network with different layer sizes/activations.

```
In [4]: plot_data(newfig=True)
        plot_model(model)
        plt.show()
```

In []: